

Quant Research OS

量化研究操作系统

Technical Handbook & System Walkthrough

技术手册与系统详解（英中对照）

Architecture · Quant Theory · Orchestration · Engines · Workstation UI · Operations

架构 · 量化理论 · 编排 · 引擎 · 工作站界面 · 运维

Hard rule: the planner/agent may propose research directions, but **deterministic quantitative engines are the sole source of truth** for financial metrics, positions, and risk numbers.

硬性规则：规划器/智能体可以提出研究方向，但确定性量化引擎是财务指标、持仓与风险数字的唯一真相来源。智能体不得编造夏普比率等指标。

Audience: quantitative researchers, platform engineers, and reviewers who need to understand how research questions become experiments, alphas, portfolio impact, and auditable reports. Layout: **English left / Chinese right**.

读者：量化研究员、平台工程师与评审者——需要理解研究问题如何变成实验、Alpha、组合影响与可审计报告。版式：左侧英文 / 右侧中文。

Chinese font embedded for this build: **SongtiSC**. Technical identifiers (API paths, IDs, code) remain in English by design.

本版嵌入中文字体：SongtiSC。技术标识符（API 路径、ID、代码）按惯例保留英文。

1. Contents

1. 目录

- 2. Product intent and design principles
- 3. Repository map
- 4. End-to-end architecture
- 5. Domain model
- 6. Research orchestrator (state machine)
- 7. Agents and tool allowlist
- 8. Cross-sectional engine and anti look-ahead
- 9. Costs, metrics, walk-forward, robustness, regimes
- 10. Alpha library, diversification, portfolio & risk
- 11. Paper trading honesty model
- 12. Storage, reproducibility, lineage
- 13. FastAPI surface and realtime events
- 14. Workstation UI architecture
- 15. Flagship FX workflow walkthrough
- 16. Testing, operations, known gaps
- 17. Appendix: key code references
- 18. Deep dive: metrics mathematics
- 19. Deep dive: backtest loop internals
- 20. Deep dive: cost model economics
- 21. Deep dive: orchestrator experiment loop
- 22. Deep dive: hypothesis catalog & falsification
- 23. Deep dive: storage schema & traces
- 24. Deep dive: API contracts & UI binding
- 25. Deep dive: traceability drill-down example
- 26. Threats to validity & research hygiene
- 27. Deep dive: walk-forward validation
- 28. Deep dive: robustness & fragility
- 29. Deep dive: regime analysis
- 30. Deep dive: diversification mathematics
- 31. Deep dive: ToolRouter & allowlist
- 32. Worked example: from returns to decision
- 33. Workstation page-by-page guide
- 34. Translation & review notes

- 2. 产品定位与设计原则
- 3. 仓库结构
- 4. 端到端架构
- 5. 领域模型
- 6. 研究编排器（状态机）
- 7. 智能体与工具白名单
- 8. 横截面引擎与反前视
- 9. 成本、指标、滚动验证、稳健性、市场状态
- 10. Alpha 库、分散化、组合与风险
- 11. 模拟盘诚实模型
- 12. 存储、可复现与血缘
- 13. FastAPI 接口与实时事件
- 14. 工作站 UI 架构
- 15. 旗舰外汇流程走读
- 16. 测试、运维与已知缺口
- 17. 附录：关键代码索引
- 18. 深入：指标数学
- 19. 深入：回测循环内部
- 20. 深入：成本模型经济学
- 21. 深入：编排器实验循环
- 22. 深入：假设目录与证伪
- 23. 深入：存储模式与轨迹
- 24. 深入：API 契约与 UI 绑定
- 25. 深入：可追溯下钻示例
- 26. 效度威胁与研究卫生
- 27. 深入：滚动前瞻验证
- 28. 深入：稳健性与脆弱性
- 29. 深入：市场状态分析
- 30. 深入：分散化数学
- 31. 深入：ToolRouter 与白名单
- 32. 演算示例：从收益到决策
- 33. 工作站逐页指南
- 34. 翻译与审阅说明

2. Product Intent and Design Principles

2. 产品定位与设计原则

Quant Research OS is an autonomous AI quantitative research laboratory. It is not a generic chat demo and not a retail trading app. The primary user is a professional researcher who must answer: what is being researched, why, which experiments ran, what survived adversarial review, how candidates interact with the existing book, and what risks remain.

Quant Research OS 是一套自主式 AI 量化研究实验室。它不是通用聊天演示，也不是零售交易 App。主要用户是专业研究员，需要回答：系统在研究什么、为什么研究、跑了哪些实验、哪些想法通过了对抗性审查、候选策略与现有组合如何相互作用、还剩哪些风险。

2.1 Design principles

2.1 设计原则

- **Precision & trust** — numbers come from engines with provenance IDs.
- **Transparency** — every transition is traced; reports link to experiments.
- **Controllability** — budgets, cooperative cancel, explicit state machine.
- **Information density** — workstation UI optimized for hours of daily use.
- **No metric invention** — agents may narrate; they must not invent Sharpe.

- 精确与信任 — 数字来自带溯源 ID 的引擎。
- 透明 — 每次状态转换可追踪；报告链接到实验。
- 可控 — 预算、协作式取消、显式状态机。
- 信息密度 — 工作站界面为每日长时间使用优化。
- 禁止编造指标 — 智能体可以叙述，但不得编造夏普比率。

2.2 Architectural hard boundary

2.2 架构硬边界

Language models are strong at planning, critique, and synthesis, but unreliable as calculators for portfolio mathematics. Therefore the system splits responsibilities.

语言模型擅长规划、批评与综合，但作为组合数学计算器并不可靠。因此系统明确拆分职责。

LLM / planner / orchestrator → proposes hypotheses, chooses tools, writes narrative
Deterministic quant engine → backtests, metrics, risk, correlations, gates
/ →
→

3. Repository Map

3. 仓库结构

Path Path	Role Role
src/.../domain/ 领域模型目录	Pydantic domain objects Pydantic 领域对象
src/.../engine/ 引擎目录	Backtest, metrics, WF, robustness... 回测、指标、滚动验证、稳健性等
src/.../agents/ 智能体目录	Deterministic planner/reviewer 确定性规划/审查 (+未来 LLM 提示词)
src/.../tools/ 工具目录	Allowlisted tool router 白名单工具路由
src/.../orchestration/ 编排目录	State machine + budgets 状态机 + 预算
src/.../alpha/ Alpha 目录	Alpha / strategy registry Alpha / 策略注册表
src/.../experiments/ 实验目录	Experiment registry + hashing 实验注册与配置哈希
src/.../storage/ 存储目录	SQLite, paths, traces SQLite、路径、轨迹
src/.../api/ API 目录	FastAPI FastAPI 接口
web/ 前端目录	Next.js workstation UI Next.js 研究工作站

tests/ 测试目录	Invariant + integration tests 不变量与集成测试
docs/ 文档目录	Architecture, audit, this handbook 架构、审计与本手册

4. End-to-End Architecture

4. 端到端架构

A research question enters through CLI, API, or the workstation command bar. The orchestrator creates a ResearchRequest, plans hypotheses, validates data, runs budgeted experiments through the tool layer, validates statistically, stresses robustness and regimes, checks diversification versus existing alphas, runs adversarial review, evaluates portfolio/risk impact, and emits a report with traces.

研究问题可通过 CLI、API 或工作站命令栏进入。编排器创建 ResearchRequest，规划假设，校验数据，经工具层在预算内运行实验，做统计验证、稳健性与市场状态压力测试，检查相对既有 Alpha 的分散性，进行对抗性审查，评估组合/风险影响，并输出带轨迹的报告。

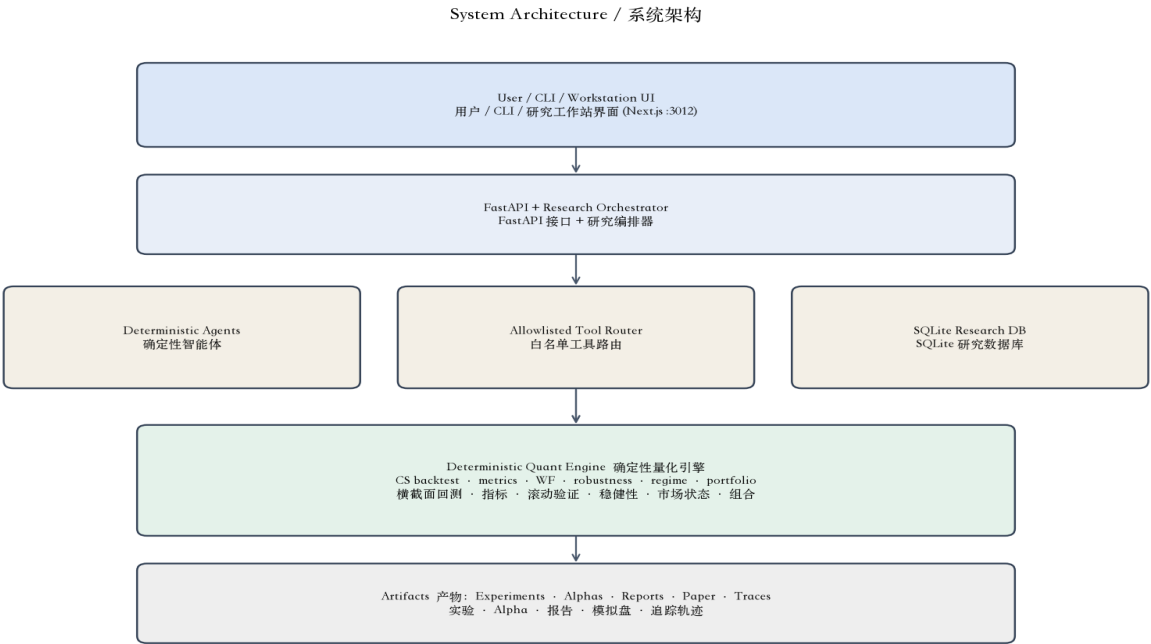


Figure 1. Layered system architecture.
图 1. 分层系统架构。

5. Domain Model

5. 领域模型

Domain objects are typed Pydantic models. They form the contract between storage, engines, API, and UI.

领域对象是带类型的 Pydantic 模型，构成存储、引擎、API 与 UI 之间的契约。

- **ResearchRequest** — user question, universe, budgets, status lifecycle.
- **ResearchPlan** — interpreted objectives, constraints, hypotheses.
- **Experiment** — strategy + parameters + dataset + status + metric linkage.
- **BacktestResult** — equity/returns artifact + performance metrics.
- **Strategy / Alpha** — reusable definition vs promoted research candidate.
- **ReviewResult** — adversarial findings with severity.
- **ResearchReport** — executive summary + sections + survivors/rejects.

- ResearchRequest** — 用户问题、标的宇宙、预算、状态生命周期。
- ResearchPlan** — 系统解释后的目标、约束与假设。
- Experiment** — 策略 + 参数 + 数据集 + 状态 + 指标关联。
- BacktestResult** — 权益/收益产物 + 绩效指标。
- Strategy / Alpha** — 可复用定义 vs 晋升后的研究候选。
- ReviewResult** — 带严重级别的对抗性发现。
- ResearchReport** — 执行摘要 + 章节 + 幸存者/否决项。

5.1 Alpha lifecycle

5.1 Alpha 生命周期

CANDIDATE → PROMISING → ROBUST → PAPER_TRADING → (LIVE future)
 ■ REJECTED / RETIRED
 # ■■ → ■■■ → ■■ → ■■■ →■■■■■
 # ■■■ / ■■

Promotion requires **metrics_source_ids** so every Sharpe/drawdown claim is traceable to a concrete backtest/experiment artifact. This is an anti-hallucination control.

晋升要求提供 **metrics_source_ids**，使每条夏普/回撤声明都能追溯到具体回测/实验产物。这是反幻觉（防编造数字）控制。

6. Research Orchestrator

6. 研究编排器

The orchestrator is an explicit finite state machine (not an unbounded agent loop). Transitions are logged to the trace table. Budgets limit hypotheses/experiments. Cancel is cooperative: status becomes CANCELLED and is checked between hypotheses.

编排器是显式有限状态机（不是无界智能体循环）。转换写入轨迹表。预算限制假设/实验数量。取消是协作式的：状态设为 CANCELLED，并在假设之间检查。

Research Orchestrator State Machine / 研究编排状态机




Figure 2. Orchestrator nodes from START to END.

图 2. 编排器节点：从 START 到 END。

6.1 Why a state machine?

6.1 为什么用状态机？

Autonomous research systems fail when agents recurse forever, skip validation, or mutate goals silently. An explicit transition table makes the laboratory auditable.

当智能体无限递归、跳过验证或静默改变目标时，自主研究系统会失败。显式转换表使实验室可审计。

```

TRANSITIONS = [
  START, ResearchPlanner, DataDiscovery, HypothesisGeneration,
  ExperimentDesign, ExperimentExecution, StatisticalValidation,
  RobustnessTesting, RegimeAnalysis, DiversificationAnalysis,
  AdversarialReview, PortfolioAnalysis, RiskReview,
  PaperTradingGate, Report, END,
]
```

7. Agents and Tool Allowlist

7. 智能体与工具白名单

Current production mode uses **deterministic agents** (agents/core.py). agents/prompts.py reserves future LLM mode. The UI System Health panel labels this honestly as “LLM services: deterministic agents”.

当前生产模式使用确定性智能体 (agents/core.py)。agents/prompts.py 预留未来 LLM 模式。UI 系统健康面板如实标注为“LLM 服务：确定性智能体”。

7.1 Planner

7.1 规划器

`plan_research()` expands a question into a fixed hypothesis set for the flagship FX workflow (momentum, reversal, carry-like structures, low-correlation constraints). It is a template, not an LLM call.

`plan_research()` 将问题展开为旗舰外汇流程的固定假设集合（动量、反转、类套利结构、低相关约束）。它是模板，不是 LLM 调用。

7.2 Tool router

7.2 工具路由

All quantitative side effects go through ToolRouter with an allowlist. Agents cannot call arbitrary Python. Typical tools: list/inspect/validate datasets, run_backtest, walk-forward, robustness batteries, correlation versus existing library, stress tests.

所有量化副作用都经过带白名单的

ToolRouter。智能体不能调用任意 Python。典型工具：列出/检查/校验数据集、run_backtest、滚动验证、稳健性套件、相对既有库的相关性、压力测试。

8. Cross-Sectional Engine & Anti Look-Ahead

8. 横截面引擎与反前视

Cross-sectional strategies rank a universe at each rebalance and hold a long/short basket. In FX G10, a classic example is momentum: buy recent winners, sell losers. The scientific risk is **look-ahead bias** — using information not available at decision time.

横截面策略在每次再平衡时对一篮子标的排序，并持有多空组合。在 FX G10 中，经典例子是动量：买入近期赢家、卖出输家。科学风险是前视偏差（look-ahead bias）——使用决策时尚不可得的信息。


8.1 Timing contract

8.1 时序契约

- At bar t close, compute signal from prices available at t .
- Schedule target weights for execution at $t + \text{execution_lag}$.
- Default $\text{execution_lag} \geq 1 \rightarrow$ no same-bar return on new weights.
- Between rebalances, share holdings are fixed; weights drift with prices.
- Turnover is measured in weight space after NAV renormalization.

在 K 线 t 收盘，仅用 t 时刻可得价格计算信号。
将目标权重安排在 $t + \text{execution_lag}$ 执行。
默认 $\text{execution_lag} \geq 1 \rightarrow$ 新权重不得计入同一根 K 线收益。
再平衡之间固定持股数量；权重随价格漂移。
换手率在净值再归一化后的权重空间中计量。

Cross-Sectional Execution Timing / 横截面执行时序



Anti look-ahead: no same-bar PnL on newly decided weights
反前视：新权重不得在同一根K线计入收益

Figure 3. Signal $\rightarrow$ lag $\rightarrow$ execution $\rightarrow$ drift.
图 3. 信号 $\rightarrow$ 滞后 $\rightarrow$ 执行 $\rightarrow$ 漂移。

8.2 Selection methods

8.2 选股/选汇方法

Supported constructions include top/bottom-N, percentile baskets, z-score weights, and rank weights. Signals include momentum and short-term reversal, plus feature generators from features/library.py (carry, volatility rank, value proxy, liquidity stability).

支持 top/bottom-N、分位数篮子、z-score 加权与秩加权。信号包括动量与短期反转，以及 features/library.py 特征生成器（carry、波动排序、价值代理、流动性稳定）。

8.3 Weight-space turnover & NAV renormalize

8.3 权重空间换手与净值再归一

- Target weights are decided on rebalance bars, then shifted by execution_lag .
- One-way turnover $= 0.5 \cdot \sum |w_{\text{new}} - w_{\text{old}}|$ (weight space, not share counts).
- Gross portfolio return uses beginning-of-bar holdings $\cdot$ asset returns.

目标权重在再平衡日决定，再按 execution_lag 平移。

单边换手 $= 0.5 \cdot \sum |w_{\text{new}} - w_{\text{old}}|$ （权重空间，非股数）。

组合毛收益 = 期初持仓 $\cdot$ 资产收益。

- Holdings drift with asset returns, then divide by $(1 + \text{port_ret})$ so next weights stay NAV-relative.
- Net returns = portfolio returns – cost series from turnover $\times$ cost rate.

If execution_lag were allowed to be 0, same-bar PnL on new weights would silently invent edge. The engine raises if lag < 1.

持仓随资产收益漂移，再除以 $(1 + \text{port_ret})$ ，使次日权重相对净值。

净收益 = 组合收益 由换手 $\times$ 成本率得到的成本序列。

若允许

execution_lag=0，新权重同日盈亏会静默伪造优势。引擎在 lag < 1 时抛错。

9. Costs, Metrics, Validation Engines

9. 成本、指标与验证引擎

9.1 Transaction costs

9.1 交易成本

Costs are first-class in engine/costs.py. Three FX presets — OPTIMISTIC (0.85 bps), BASELINE (2.75 bps), PESSIMISTIC (10 bps) variable — multiply one-way turnover. Strategies must remain attractive under BASELINE; promotion narratives require survival under PESSIMISTIC stress. See deep dive §20 for the full economics table.

成本在 engine/costs.py 中是一等公民。三种外汇预设 — OPTIMISTIC (0.85 bps)、BASELINE (2.75 bps)、PESSIMISTIC (10 bps) 可变 — 乘以单边换手。策略在 BASELINE 下仍需有吸引力；晋升叙述需经受 PESSIMISTIC 压力。完整经济学表见深入 §20。

9.2 Performance metrics

9.2 绩效指标

Metrics are computed exclusively in engine/metrics.py with pinned definitions (sibling libraries disabled — sibling_metrics_available() → False). Critical pin: population std (ddof=0). Sortino uses full-sample downside deviation $\sqrt{\text{mean}(\min(r, 0)^2)}$, not std of negatives only. Deep dive §18 lists the exact formulas and code contract.

指标仅在 engine/metrics.py 按固定定义计算（兄弟库禁用 — sibling_metrics_available() → False）。关键固定：总体标准差 (ddof=0)。索提诺使用全样本下行偏差 $\sqrt{\text{mean}(\min(r, 0)^2)}$ ，而非仅负收益子集标准差。深入 §18 列出精确公式与代码契约。

- **Annual return / volatility** — 252-day annualization.
- **Sharpe** — mean excess / std(ddof=0) · $\sqrt{252}$.
- **Sortino** — full-sample downside deviation.
- **Max drawdown / Calmar** — path-dependent risk.
- **Turnover / costs / trade_count** — implementability.

年化收益 / 波动率 — 252 日年化。

夏普 — 超额均值 / std(ddof=0) · $\sqrt{252}$ 。

索提诺 — 全样本下行偏差。

最大回撤 / 卡玛 — 路径依赖风险。

换手 / 成本 / 成交笔数 — 可实施性。

Performance Measurement Intuition / 绩效度量直觉




Figure 4. Intuition for equity path and rolling Sharpe.

图 4. 权益路径与滚动夏普的直觉示意。

9.3 Walk-forward

9.3 滚动前瞻验证 (Walk-forward)

Walk-forward (engine/walk_forward.py) evaluates a **frozen** cross-sectional rule on expanding or rolling windows (defaults: 126 train / 63 test / step 63). Parameters are not re-optimized inside train windows. OOS segments are de-duplicated before aggregate Sharpe. OOS — not in-sample — is the primary promotion signal. Deep dive §27.

Walk-forward (engine/walk_forward.py) 在扩展或滚动窗上评估冻结横截面规则（默认：126 训练 / 63 测试 / 步长 63）。训练窗内不重新优化参数。汇总夏普前对样本外片段去重。晋升看 OOS 而非样本内。深入 § 27。

9.4 Robustness

9.4 稳健性

Robustness (engine/robustness.py) sweeps lookback (and similar) grids and flags **fragile** knife-edge peaks (interior peak Sharpe – neighbor mean > 0.5 with peak > 0.3) and jagged surfaces (mean |ΔSharpe| ≥ 0.35). Adversarial review treats fragility as a promotion blocker. Deep dive §28.

稳健性 (engine/robustness.py) 扫描 lookback (等) 网格，标记脆弱刀锋峰（内部峰值夏普 – 邻域均值 > 0.5 且峰值 > 0.3）与锯齿曲面（平均 |Δ夏普| ≥ 0.35）。对抗性审查将脆弱性视为晋升阻断。深入 § 28。

9.5 Regimes

9.5 市场状态 (Regime)

Regime labels (engine/regime.py) combine rolling-vol tertiles with risk-on/off trend, then shift by 1 bar to avoid same-day attribution leakage. Labels are heuristic (confidence ~0.4). Concentration across a single regime slice is flagged. Deep dive §29.

市场状态标签 (engine/regime.py) 结合滚动波动三分位与风险偏好趋势，再平移 1 根 K 线以避免同日归因泄漏。标签为启发式（置信度 ~0.4）。绩效集中在单一状态会被标记。深入 § 29。

10. Alpha Library, Diversification, Portfolio & Risk

10. Alpha 库、分散化、组合与风险

The alpha library stores candidates with status, metrics, and correlation to the existing book. A seeded momentum alpha (ALP-existing-momentum) represents incumbent exposure for low-correlation discovery.

Alpha 库保存候选的状态、指标及与现有组合的相关性。种子动量 Alpha (ALP-existing-momentum) 代表既有敞口，用于低相关发现。

10.1 Diversification analysis

10.1 分散化分析

Return correlations, downside correlations, and incremental Sharpe hints estimate whether a candidate is genuine diversification or a duplicate risk bet.

收益相关、下行相关与增量夏普提示，用于判断候选是真正分散，还是重复风险押注。

10.2 Portfolio / risk

10.2 组合 / 风险

The workstation Portfolio page supports what-if: compare current book vs book + Alpha X on Sharpe, correlation, and counts. Risk Center highlights concentration and high-correlation alerts.

工作站组合页支持情景分析：比较当前组合 vs 组合 + Alpha X 的夏普、相关性与数量。风险中心突出集中度与高相关告警。

11. Paper Trading Honesty Model

11. 模拟盘诚实模型

Paper trading in v0.2 may simulate paths using IID noise calibrated to backtest μ/σ . This is explicitly labeled in payloads and UI banners. It is useful for monitoring plumbing, but it is **not** a claim of realistic market replay.

v0.2 的模拟盘可能使用按回测 μ/σ 校准的 IID 噪声路径。这在载荷与 UI 横幅中明确标注。它对监控链路有用，但不是真实市场回放的声明。

The UI forces a clear distinction among **BACKTEST**, **PAPER TRADING**, and **LIVE**. Confusing these modes is a primary operational failure mode in quant platforms.

UI 强制区分 回测 (BACKTEST)、模拟盘 (PAPER TRADING) 与 实盘 (LIVE)。混淆这些模式是量化平台的主要运维失败模式之一。

12. Storage, Reproducibility, Lineage

12. 存储、可复现与血缘

SQLite (WAL mode, FK pragma, indexed JSON tables) stores research requests, plans, experiments, alphas, reports, paper state, checkpoints, and agent traces. QROS_DATA_ROOT controls the data directory (default data/local).

SQLite (WAL、外键、带索引 JSON 表) 存储研究请求、计划、实验、Alpha、报告、模拟盘状态、检查点与智能体轨迹。QROS_DATA_ROOT 控制数据目录 (默认 data/local)。

12.1 Config hashing

12.1 配置哈希

Experiment configs are hashed so identical parameterizations can be detected and reproduced.

实验配置会被哈希，以便发现并复现相同参数化。

Traceability Lineage / 可追溯血缘




Figure 5. Claim → evidence lineage used by reports and UI.

图 5. 报告与 UI 使用的「结论 → 证据」血缘。

13. FastAPI Surface & Realtime

13. FastAPI 接口与实时

The API exposes health, research CRUD/cancel, traces, experiments, strategies, alphas, backtests, reports, portfolio, correlations, paper stepping, and SSE /events/stream. Optional QROS_API_KEY. CORS defaults include the Next.js origin :3012. Experiment list responses are enriched with metrics joined from backtest_results.

API 提供健康检查、研究创建/查询/取消、轨迹、实验、策略、Alpha、回测、报告、组合、相关性、模拟盘步进，以及 SSE /events/stream。可选 QROS_API_KEY。CORS 默认包含 Next.js :3012。实验列表会从 backtest_results 关联填充指标。

```

GET /health
POST /research
GET /research/{id} · GET /research/{id}/trace
POST /research/{id}/cancel
GET /experiments[?research_id=] # includes joined metrics
GET /alphas · /strategies · /portfolio · /paper
GET /events/stream
  
```

14. Workstation UI Architecture

14. 工作站 UI 架构

The UI (web/) is a Next.js App Router application with a persistent shell: collapsible left navigation, top status bar, main canvas, and a bottom research command bar. Design tokens encode a dark research-terminal aesthetic (IBM Plex). Server state uses SWR.

UI (web/) 是 Next.js App Router 应用，具有持久壳层：可折叠左侧导航、顶部状态栏、主画布与底部研究命令栏。设计令牌呈现深色研究终端美学（IBM Plex）。服务端状态使用 SWR。

Workstation Shell / 工作站壳层布局



Figure 6. Application shell composition.
图 6. 应用壳层组成。

14.1 Screen inventory

14.1 页面清单

Route ■■ Route	Purpose ■■ Purpose
/	Overview dashboard 总览仪表盘
/research/[id] /research/[id]	Flagship research workspace 旗舰研究工作区
/experiments /experiments	Experiment ledger + export 实验台账 + 导出
/alphas/[id] /alphas/[id]	Alpha detail + lineage Alpha 详情 + 血缘
/portfolio /portfolio	Allocation + what-if 配置 + 情景分析
/risk /risk	Risk center 风险中心
/agents /agents	Agent observability 智能体可观测性
/reports/[id] /reports/[id]	Report reader 报告阅读器
/memory /memory	Research memory search 研究记忆搜索
/paper /paper	Paper trading (bannered) 模拟盘（带模式横幅）
Cmd+K palette Cmd+K 命令面板	Global object search 全局对象搜索

14.2 Frontend module boundaries

14.2 前端模块边界

web/src/app/*	pages ■■
web/src/components/shell	chrome ■■
web/src/components/ui	design system ■■■■
web/src/components/charts	TimeSeriesChart ■■■■
web/src/components/research	ResearchGraph ■■■■
web/src/domain/types.ts	typed models ■■■■
web/src/lib/api.ts	API client API ■■■■
web/src/lib/realtime.ts	SSE ■■
web/src/styles/*	tokens/CSS ■■■■

15. Flagship FX Workflow Walkthrough

15. 旗舰外汇流程走读

Flagship question: find a robust cross-sectional FX strategy with low correlation to existing momentum. The laboratory expands hypotheses, runs budgeted backtests, validates OOS, measures correlation to ALP-existing-momentum, adversarially reviews survivors, and writes a report.

旗舰问题：寻找与既有动量低相关、稳健的外汇横截面策略。实验室展开假设，在预算内回测，做样本外验证，测量与 ALP-existing-momentum 的相关性，对幸存者做对抗性审查并撰写报告。

```
quant research run \
  "Find a robust cross-sectional FX strategy with low correlation "\
  "to my existing momentum strategies." \
  --max-experiments 8

# Then open workstation Research workspace for that research_id.
# ■■■■■■■■■■ research_id ■■■■■■■■■■
```

15.1 Step-by-step laboratory path

15.1 实验室逐步路径

- Planner emits competing FX hypotheses (carry, reversal, vol, value, liquidity) with falsifiers.
- Data agent validates synthetic FX panels and seeds ALP-existing-momentum as incumbent book.
- Each hypothesis becomes a hashed Experiment → lagged CS backtest → metrics artifact.
- Walk-forward / robustness / regime filters remove fragile winners.
- Diversification vs momentum + adversarial review assign PROMISING / REJECT / more-research.
- Report + traces enable drill-down; Portfolio what-if estimates book impact.

规划器发出竞争的外汇假设（套息、反转、波动、价值、流动性）及证伪条件。

数据智能体校验合成外汇面板，并播种 ALP-existing-momentum 作为既有组合。

每个假设变成带哈希的 Experiment → 滞后横截面回测 → 指标产物。

滚动验证 / 稳健性 / 市场状态过滤去掉脆弱赢家。

相对动量的分散分析 + 对抗性审查给出 PROMISING / REJECT / 需更多研究。

报告 + 轨迹支持下钻；组合情景分析估计账面影响。

15.2 How to read a completed run

15.2 如何阅读已完成运行

In the workstation: Overview → latest research card → Research workspace graph (all nodes completed) → Experiments tab sorted by Sharpe → open best EXP → confirm costs/lag/dataset → Agent Inspector tool payloads → Report survivors → Portfolio what-if. Never stop at the executive summary.

在工作站：总览 → 最新研究卡片 → 研究工作区图（节点均完成）→ 实验页按夏普排序 → 打开最佳 EXP → 确认成本/滞后/数据集 → 智能体检查器工具载荷 → 报告幸存者 → 组合情景分析。切勿只看执行摘要。

16. Testing, Operations, Known Gaps

16. 测试、运维与已知缺口

16.1 Verification snapshot

16.1 验证快照

- Backend: pytest — 32 passed.
- Frontend: vitest + next build — passed.
- Runtime smoke: API healthy; UI routes HTTP 200; browser confirmed Overview, Research graph, Experiments with Sharpes, Alpha Library, Portfolio.

后端：pytest — 32 通过。

前端：vitest + next build — 通过。

运行冒烟：API 健康；UI 路由 HTTP 200；浏览器确认总览、研究图、带夏普的实验表、Alpha 库、组合页。

- Fixes during verification: normalize user_question; enrich experiment metrics from backtests.

验证中修复：规范化 user_question；从回测关联填充实验指标。

16.2 Runbook

16.2 运行手册

```
# API
source .venv/bin/activate
quant serve --host 127.0.0.1 --port 8002

# UI
cd web && npm install && npm run dev    # http://127.0.0.1:3012

# Handbook PDF
python docs/generate_handbook_pdf.py
```

16.3 Known gaps (honest)

16.3 已知缺口（如实）

- Planner is a deterministic template — not a live LLM.
- Paper trading may use IID noise simulation (labeled).
- Full resume-from-checkpoint node restart is incomplete.
- True async worker fleet / RBAC / Docker hardening incomplete.
- Some UI charts are metric-conditioned until full equity series attach.

规划器是确定性模板——不是实时 LLM。
模拟盘可能使用 IID 噪声仿真（已标注）。
从检查点按节点完整恢复尚未完成。
真正的异步工作集群 / RBAC / Docker 加固未完成。
部分 UI 图表在完整权益序列接入前由指标条件生成示意。

17. Appendix — Key Code References

17. 附录 — 关键代码索引

orchestration/runner.py	state machine + budgets + cancel █████/██/██
agents/core.py	plan_research, adversarial_review ██████████
tools/router.py	allowlisted tools ██████
engine/cross_sectional.py	lagged CS backtest ██████████
engine/metrics.py	pinned performance math ██████████
engine/walk_forward.py	OOS evaluation ██████
engine/robustness.py	sensitivity batteries ██████
engine/regime.py	shifted regime labels ██████████
alpha/registry.py	alpha/strategy lifecycle ██████
experiments/registry.py	experiment persistence ██████
storage/db.py	SQLite + traces + checkpoints
api/app.py	HTTP + SSE + enrichment
web/src/app/research/[id]	flagship workspace UI ██████
web/src/lib/api.ts	typed client ██████████

18. Deep Dive: Metrics Mathematics

18. 深入解析：指标数学

All performance numbers in Quant Research OS are computed exclusively in engine/metrics.py. Sibling metric libraries are **disabled** (sibling_metrics_available() → False) because different ddof conventions silently change Sharpe. The pinned formulas below are the contract — agents and tools must not invent numbers.

Quant Research OS 中的全部绩效数字仅在 engine/metrics.py 计算。兄弟指标库被禁用 (sibling_metrics_available() → False)，因为不同的 ddof 约定会静默改变夏普。以下固定公式即为契约——智能体与工具不得编造数字。

18.1 Equity path and annualization

18.1 权益路径与年化

Given net daily returns r_t , equity is $E_t = \prod(1 + r_t)$. Cumulative return is $E_T - 1$. Annualized return uses $(1 + \text{cum})^{(252/N)} - 1$ with N = number of daily observations and 252 trading days per year.

给定净日收益 r_t ，权益为 $E_t = \prod(1 + r_t)$ 。累计收益为 $E_T - 1$ 。年化收益使用 $(1 + \text{cum})^{(252/N)} - 1$ ， N 为日观测数，每年 252 个交易日。

18.2 Sharpe, Sortino, MDD (pinned)

18.2 夏普、索提诺、最大回撤（固定）

• **Volatility**: $\sigma = \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$. Population standard deviation (ddof=0) is pinned — not sample std.

波动率: $\sigma = \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$ 。固定使用总体标准差 (ddof=0) ——非样本标准差。

• **Sharpe**: $\text{mean}(r - r_f/252) / \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$. Excess return over daily risk-free; denominator uses full-sample std with ddof=0.

夏普: $\text{mean}(r - r_f/252) / \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$ 。超额为日无风险利率之上；分母为全样本 ddof=0 标准差。

• **Sortino**: downside deviation = $\sqrt{\text{mean}(\min(r, 0)^2)}$ over the full sample (zero returns contribute). Not std of negative returns only.

索提诺: 下行偏差 = 全样本上的 $\sqrt{\text{mean}(\min(r, 0)^2)}$ （零收益也计入）。不是仅对负收益子集求标准差。

• **Max drawdown (MDD)**: $\min(E_t / \text{cummax}(E) - 1)$.

最大回撤 (MDD) : $\min(E_t / \text{cummax}(E) - 1)$ 。

• **Calmar**: annualized return / |MDD| when MDD < 0.


卡玛: 当 MDD < 0 时，年化收益 / |MDD|。

```
# engine/metrics.py – pinned implementation
equity = (1 + r).cumprod()
ann_ret = (1 + cum) ** (252 / len(r)) - 1
vol = r.std(ddof=0) * sqrt(252)
excess = r - risk_free_rate / 252
sharpe = excess.mean() / r.std(ddof=0) * sqrt(252)
downside = np.minimum(r, 0.0)
down_dev = sqrt(np.mean(downside ** 2)) # full-sample
sortino = excess.mean() / down_dev * sqrt(252)
mdd = (equity / equity.cummax() - 1.0).min()
calmar = ann_ret / abs(mdd) if mdd < 0 else 0.0
```

If a reported Sharpe cannot be reproduced by these formulas on the linked return series, the number is invalid — regardless of narrative quality.

若报告的夏普无法用这些公式在关联收益序列上复现，则该数字无效——无论叙述多么有说服力。

Metrics Mathematics Pipeline / 指标数学流水线



ddof=0 pinned · full-sample Sortino downside · no sibling imports
固定 ddof=0 · 全样本索提诺下行 · 不导入兄弟库

Figure 7. Pinned metrics pipeline from net returns to report fields.

图 7. 从净收益到报告字段的固定指标流水线。

19. Deep Dive: Backtest Loop Internals

19. 深入解析：回测循环内部

File: engine/cross_sectional.py. The cross-sectional engine constructs long/short baskets from scores and simulates portfolio returns under a strict anti-look-ahead timing contract: signal at bar t , execution at $t + \text{execution_lag}$ (≥ 1), weight-space turnover, NAV renormalization, and transaction costs.

文件: engine/cross_sectional.py. 横截面引擎由分数构建多空篮子, 并在严格反前视时序契约下模拟组合收益: t 时刻信号、 $t + \text{execution_lag}$ (≥ 1) 执行、权重空间换手、净值再归一化与交易成本。

19.1 Configuration surface

19.1 配置面

```
class CrossSectionalConfig(BaseModel):
    lookback: int = 20
    top_n: int = 3
    bottom_n: int = 3
    selection: TOP_BOTTOM_N | PERCENTILE | ZSCORE_WEIGHT | RANK_WEIGHT
    rebalance_every: int = 5
    execution_lag: int = Field(default=1, ge=1) # hard floor
    gross_exposure: float = 1.0
    cost_assumption: optimistic | baseline | pessimistic
    signal_name: str = "momentum" # or reversal / feature library
```

19.2 Selection methods

19.2 选择方法

Method 方法	Mechanism 机制	When to use 适用场景
TOP_BOTTOM_N Top/Bottom N	Long top N, short bottom N with equal split 多头 top N、空头 bottom N 等权	Default FX basket; interpretable book 默认外汇篮子; 可解释持仓
PERCENTILE 分位数	Long above percentile_long; short below percentile_short 高于上分位做多; 低于下分位做空	Smooth basket sizing across universe 跨宇宙平滑篮子规模
ZSCORE_WEIGHT z 分数加权	Center scores; weight proportional to z-score 中心化分数; 按 z 分数比例加权	Continuous exposure to signal strength 对信号强度连续暴露
RANK_WEIGHT 秩加权	Center ranks; weight proportional to rank 中心化秩; 按秩比例加权	Robust to outliers in raw scores 对原始分数异常值稳健

19.3 Simulation loop (critical)

19.3 仿真循环（关键）

- On rebalance bars, compute target weights from cross-sectional scores.
- Shift targets by `execution_lag` bars (≥ 1). Default `lag=1` $\rightarrow$ no same-bar PnL on newly decided weights.
- Turnover in **weight space**: $\text{turnover} = 0.5 \cdot \sum |w_{\text{new}} - w_{\text{old}}|$. One-way traded fraction.
- Portfolio return: $\text{port_ret} = \sum (w \cdot r_{\text{asset}})$ using beginning-of-bar weights.
- NAV renormalize**: drift holdings by asset returns, then divide by $(1 + \text{port_ret})$ so next-day weights remain return-relative.
- Subtract cost series derived from $\text{turnover} \times \text{cost rate}$.

```
# Anti look-ahead guard
if cfg.execution_lag < 1:
    raise ValueError("execution_lag must be >= 1")
scheduled = target.shift(cfg.execution_lag)
turnover = abs(new_w - cur_w).sum() / 2.0 # weight space
port_ret = (holdings * asset_rets).sum()
holdings = holdings * (1 + asset_rets) / (1 + port_ret) # NAV renormalize
net = ret_series - cost_series
```

在再平衡日由横截面分数计算目标权重。

将目标按 `execution_lag` 根 K 线平移 (≥ 1)。默认 `lag=1` $\rightarrow$ 新权重不得在同一根 K 线计入收益。

权重空间换手: $\text{turnover} = 0.5 \cdot \sum |w_{\text{new}} - w_{\text{old}}|$ 。单边交易比例。

组合收益: $\text{port_ret} = \sum (w \cdot r_{\text{asset}})$ ，使用期初权重。

净值再归一：按资产收益漂移持仓，再除以 $(1 + \text{port_ret})$ ，使次日权重仍相对净值。

减去由换手 $\times$ 成本率导出的成本序列。

Cross-Sectional Backtest Loop / 横截面回测循环




Figure 8. One bar of the cross-sectional simulation loop.

图 8. 横截面仿真循环中的一根 K 线。

20. Deep Dive: Cost Model Economics

20. 深入解析：成本模型经济学

File: `engine/costs.py`. Transaction costs are applied to one-way turnover ($0.5 \cdot \sum |\Delta w|$). Three FX-oriented presets bracket realistic execution: OPTIMISTIC, BASELINE, and PESSIMISTIC. Variable cost rate = (proportional + spread + slippage) bps / 10,000.

文件: `engine/costs.py`。交易成本作用于单边换手 ($0.5 \cdot \sum |\Delta w|$)。三种外汇取向预设框定现实执行: OPTIMISTIC (乐观)、BASELINE (基准)、PESSIMISTIC (悲观)。可变成本率 = (比例 + 点差 + 滑点) bps / 10,000。

Assumption 假设	Prop bps 比例 bps	Spread bps 点差 bps	Slippage bps 滑点 bps	Total var bps 可变合计 bps
OPTIMISTIC 乐观	0.5 0.5	0.25 0.25	0.1 0.1	0.85 0.85

BASELINE 基准	1.5 1.5	0.75 0.75	0.5 0.5	2.75 2.75
PESSIMISTIC 悲观	5.0 5.0	3.0 3.0	2.0 2.0	10.0 10.0
Cost per rebalance $\approx$ turnover $\times$ variable_bps / 10,000. Strategies must remain attractive under BASELINE costs; promotion pressure requires survival under PESSIMISTIC assumptions. Walk-forward and robustness batteries re-run with cost stress.			每次再平衡成本 $\approx$ turnover $\times$ variable_bps / 10,000。策略在 BASELINE 成本下仍需有吸引力；晋升压力要求经受 PESSIMISTIC 假设。滚动验证与稳健性套件以成本压力重跑。	
A strategy with Sharpe 2.0 under OPTIMISTIC but negative Sharpe under PESSIMISTIC is a cost-sensitivity failure, not a candidate for promotion.			在 OPTIMISTIC 下夏普 2.0、在 PESSIMISTIC 下夏普为负的策略是成本敏感失败，不是晋升候选。	

21. Deep Dive: Orchestrator Experiment Loop

21. 深入解析：编排器实验循环

File: orchestration/runner.py. After plan_research() and dataset validation, the runner iterates candidate hypotheses. Each iteration is budget-gated (budget.can_run_experiment()) and cancel-aware (db.is_cancelled(research_id) checked between hypotheses).

文件: orchestration/runner.py。在 plan_research() 与数据集校验之后，运行器遍历候选假设。每次迭代受预算门控 (budget.can_run_experiment()) 并感知取消 (假设间检查 db.is_cancelled(research_id))。

21.1 Feature → dataset mapping

21.1 特征 → 数据集映射

Feature 特征	SIGNAL_MAP key 映射键	Dataset 数据集	Lookback (typical) 回看 (典型)
carry carry	cross_sectional_carry cross_sectional_carry	fx_synthetic_momentum fx_synthetic_momentum	60 60
reversal reversal	short_term_reversal short_term_reversal	fx_synthetic_meanrev fx_synthetic_meanrev	1 1
volatility volatility	volatility_ranked volatility_ranked	fx_synthetic_momentum fx_synthetic_momentum	20 20
value value	value_proxy value_proxy	fx_synthetic_meanrev fx_synthetic_meanrev	20 20
liquidity liquidity	liquidity_stability liquidity_stability	fx_synthetic_momentum fx_synthetic_momentum	20 20

21.2 Per-hypothesis pipeline

21.2 每个假设的流水线

- Map hypothesis name → feature via SIGNAL_MAP.
- Persist Strategy with lookback, signal, rebalance_every; link hypothesis edge.
- Create Experiment; call run_backtest via ToolRouter.
- Run walk-forward, robustness, regime, diversification vs seeded momentum alpha.
- Adversarial review emits severity-tagged findings; survivors may register as Alpha.

经 SIGNAL_MAP 将假设名映射到特征。

持久化 Strategy (含 lookback、signal、rebalance_every)；链接假设边。


创建 Experiment；经 ToolRouter 调用 run_backtest。

运行滚动验证、稳健性、市场状态、相对种子动量 Alpha 的分散性分析。

对抗性审查产出带严重级别的发现；幸存者可能注册为 Alpha。

```
# Pseudocode – orchestration/runner.py inner loop
for hyp in plan.candidate_hypotheses:
    if db.is_cancelled(research_id): break
    if not budget.can_run_experiment(): break
    feature = SIGNAL_MAP.get(hyp.name, "momentum")
    dataset = DATASET_FOR[feature] # momentum vs meanrev panel
    strategy = Strategy(..., parameters={lookback, signal, rebalance_every})
    budget.consume_experiment()
    exp = tools.call("create_experiment", hypothesis_id=hyp.hypothesis_id, ...)
    bt = tools.call("run_backtest", experiment_id=exp.id, dataset_id=dataset, ...)
    wf = tools.call("walk_forward", ...)
    rob = tools.call("robustness_battery", ...)
    div = tools.call("analyze_diversification", vs="ALP-existing-momentum")
    rev = adversarial_review(candidate, evidence)
    maybe_register_alpha(bt, metrics_source_ids=[bt.backtest_id])
```

Orchestrator Experiment Pipeline / 编排器实验流水线



Cancel checked between hypotheses · fx_synthetic_momentum / fx_synthetic_meanrev
假设间检查取消 · 动量盘 / 均值回复盘数据集

Figure 9. Budgeted hypothesis → experiment → validation loop.
图 9. 有预算的 假设→实验→验证 循环。

22. Deep Dive: Hypothesis Catalog & Falsification

22. 深入解析：假设目录与证伪

File: agents/core.py — plan_research() emits five fixed candidate hypotheses. Each carries explicit falsification_criteria so researchers know in advance what evidence would reject the idea.

文件: agents/core.py — plan_research() 产出五个固定候选假设。每个假设带有明确的 falsification_criteria，使研究员事先知道何种证据会否决该想法。

Hypothesis 假设	Economic intuition 经济直觉	Falsification criteria 证伪标准
cross_sectional_carry 横截面 carry	High-yield currencies earn carry premium 高息货币赚取 carry 溢价	OOS Sharpe ≤ 0 after costs or corr(momentum) > 0.5 成本后 OOS 夏普 ≤ 0 或与动量相关 > 0.5
short_term_reversal 短期反转	Transitory flows reverse over short horizons 短期资金流反转	Fails walk-forward or excessive turnover under costs 滚动验证失败或成本下换手过高
volatility_ranked 波动率排序	Relative vol forecasts risk-adjusted opportunity 相对波动预测风险调整机会	No OOS edge; pure risk packaging without alpha 无 OOS 优势；纯风险包装无 alpha
value_proxy 价值代理	Mean-reversion to longer-run fair value 向长期公允价值均值回复	Collapses to slow momentum/reversal with high corr 坍缩为高相关的慢动量/反转
liquidity_stability 流动性稳定	Stable/liquid names earn premium in stress 稳定/高流动名称在压力下溢价	Indistinguishable from low-vol factor; no diversification 与低波因子不可区分；无分散性

Success criteria (plan-level): OOS Sharpe > 0.3 after baseline costs, |corr(momentum)| < 0.35, parameter surface not fragile, adversarial severity < CRITICAL.

成功标准（计划级）：基准成本后 OOS 夏普 > 0.3，|corr(动量)| < 0.35，参数面不脆弱，对抗性严重度 < CRITICAL。

23. Deep Dive: Storage Schema & Trace Model

23. 深入解析：存储模式与轨迹模型

File: storage/db.py. SQLite stores JSON payloads with indexed columns for status queries. WAL mode and foreign_keys=ON are enabled. Every orchestrator action writes structured agent_traces rows.

文件: storage/db.py。SQLite 存储 JSON 载荷，索引列支持状态查询。启用 WAL 与 foreign_keys=ON。每个编排动作写入结构化 agent_traces 行。

```
research_requests(research_id PK, payload, status, created_at)
research_plans(research_id PK → requests, payload, created_at)
hypotheses(hypothesis_id PK, research_id, payload)
experiments(experiment_id PK, research_id, payload, status,
             configuration_hash, created_at)
strategies(strategy_id PK, payload, version)
alphas(alpha_id PK, strategy_id, payload, status)
backtest_results(backtest_id PK, experiment_id, payload)
validation_results(validation_id PK, experiment_id, payload)
reports(research_id PK, payload, decision, created_at)
agent_traces(id, research_id, agent, event_type, payload, created_at)
lineage_edges(src_type, src_id, rel, dst_type, dst_id, meta)
checkpoints(research_id, node, payload, ...)
```

23.1 Lineage edges and traces

23.1 血缘边与轨迹

- research —has_hypothesis→ hypothesis
- hypothesis —implements→ strategy
- strategy —used_in→ experiment
- experiment —produced→ backtest
- experiment —validated_by→ validation
- Alpha.metrics_source_ids → backtest_id(s) for provenance
- agent_traces: planner, data, orchestrator, adversarial events with JSON payloads

research —has_hypothesis→ hypothesis

hypothesis —implements→ strategy

strategy —used_in→ experiment

experiment —produced→ backtest

experiment —validated_by→ validation

Alpha.metrics_source_ids → backtest_id (们) 用于溯源

agent_traces: 规划、数据、编排、对抗性事件及 JSON 载荷

24. Deep Dive: API Contracts & UI Binding

24. 深入解析：API 契约与 UI 绑定

File: api/app.py. Base URL default: http://127.0.0.1:8002. Optional header X-API-Key when QROS_API_KEY is set. CORS allows workstation origin :3012.	文件: api/app.py。默认基址: http://127.0.0.1:8002。若设置 QROS_API_KEY, 需头 X-API-Key。CORS 允许工作站来源 :3012。
--	---

24.1 Experiment enrichment

24.1 实验富化

GET /experiments returns registry rows joined with backtest_results metrics. Aliases ann_return and ann_vol are exposed so the workstation experiment table Sharpe column is never blank when a backtest exists.	GET /experiments 返回注册表行并关联 backtest_results 指标。暴露别名 ann_return 与 ann_vol, 使工作站实验表在存在回测时夏普列不为空。
--	--

24.2 SWR vs UI state

24.2 SWR 与 UI 状态

• Server state (SWR): keys like overview-research, research-{id}, experiments-all. Stale-while-revalidate polling. • UI state (React): selected graph node, table sort, sidebar collapse, command palette query — local component state / localStorage. • SSE: GET /events/stream emits research.updated. EventSource cannot attach API keys; UI falls back to SWR when key required. <pre>GET /research/{id} → { request: { research_id, user_question, budget, status, ... }, plan: { objectives, hypotheses, success_criteria, ... }, report: { executive_summary, survivors, decision, ... }, checkpoint: { node, payload } null } GET /research/{id}/trace → [{ agent, event_type, payload, created_at }, ...]</pre>	服务端状态 (SWR) : 键如 overview-research、research-{id}、experiments-all。陈旧-重验证轮询。 UI 状态 (React) : 选中图节点、表排序、侧栏折叠、命令面板查询 — 组件本地状态 / localStorage。 SSE: GET /events/stream 发出 research.updated。EventSource 无法附加 API key; 需要密钥时 UI 回退 SWR。
--	---

25. Deep Dive: Traceability Drill-Down Example

25. 深入解析：可追溯下钻示例

Worked path for a flagship FX research run (synthetic panels). Numbers vary by seed; the evidence chain structure is invariant.	旗舰外汇研究运行的走查路径（合成面板）。数字因种子而异；证据链结构不变。
---	--------------------------------------

Step 步骤	Entity / field 实体/字段	What to inspect 检查内容
1 1	Report.decision Report.decision	REQUIRES_MORE_RESEARCH — honest outcome REQUIRES_MORE_RESEARCH — 诚实结果
2 2	Report.survivors[0].alpha_id 幸存者 alpha_id	Link to alpha registry row 链接 Alpha 注册表行
3 3	Alpha.metrics_source_ids metrics_source_ids	List of backtest_id values — provenance anchor backtest_id 列表 — 溯源锚点
4 4	backtest_results.payload 回测载荷	Sharpe, Sortino, MDD, turnover, cost breakdown 夏普、索提诺、MDD、换手、成本分解
5 5	experiments.configuration_hash 配置哈希	Reproducibility fingerprint (params + dataset + costs) 可复现指纹（参数+数据+成本）
6 6	agent_traces (orchestrator) 编排轨迹	Tool call params/results for run_backtest, walk_forward run_backtest、walk_forward 工具调用参数/结果

• Report conclusion → survivor alpha_id	报告结论 → 幸存者 alpha_id
• Alpha.metrics_source_ids → backtest_id	Alpha.metrics_source_ids → backtest_id
• backtest_id → experiment_id → strategy parameters	backtest_id → experiment_id → 策略参数

- Trace → exact tool inputs that produced the number

轨迹 → 产生该数字的精确工具输入

Traceability Drill-Down Example / 可追溯下钻示例




Figure 10. Report → experiment → backtest → trace drill-down.

图 10. 报告 → 实验 → 回测 → 轨迹下钻。

26. Threats to Validity & Research Hygiene

26. 效度威胁与研究卫生

Quant Research OS is a laboratory, not a live trading system. Researchers must actively guard against threats that make backtests look better than deployable reality.

Quant Research OS 是实验室，不是实盘系统。研究员必须主动防范使回测看起来优于可部署现实的威胁。

26.1 Threat catalog

26.1 威胁目录

Threat 威胁	Manifestation in QROS 在 QROS 中的表现	Mitigation 缓解
Look-ahead bias 前视偏差	Same-bar execution on new weights 新权重同一根 K 线执行	execution_lag ≥ 1 enforced in engine 引擎强制 execution_lag ≥ 1
Overfitting / data snooping 过拟合/数据窥探	Many hypotheses on same panel 同一面板上多假设	Budget caps, walk-forward, adversarial review 预算上限、滚动验证、对抗性审查
Cost underestimation 成本低估	Ignoring spread/slippage 忽略点差/滑点	Three-tier FX cost presets; pessimistic stress 三层 FX 成本预设；悲观压力
Metric inconsistency 指标不一致	ddof=1 Sharpe from external lib 外部库 ddof=1 夏普	Pinned metrics.py; sibling imports disabled 固定 metrics.py；禁用兄弟库
Narrative invention 叙述编造	Agent prose without tool backing 无工具支撑的智能体 prose	Numbers only via ToolRouter; metrics_source_ids 数字仅经 ToolRouter；metrics_source_ids
Mode confusion 模式混淆	PAPER / synthetic treated as LIVE 模拟盘/合成数据当成实盘	UI banners; BACKTEST labels on synthetic data UI 横幅；合成数据 BACKTEST 标签
Survivorship in alpha library Alpha 库幸存者偏差	Only winners visible 仅见赢家	Failed experiments persisted; adversarial findings surfaced 失败实验持久化；对抗性发现展示

26.2 Practitioner checklist

26.2 从业者清单

- Verify Sharpe by recomputing from stored net returns with ddof=0 formulas.

用 ddof=0 公式从已存净收益重算以验证夏普。

- Check OOS walk-forward Sharpe, not just in-sample.
- Re-run under PESSIMISTIC costs before any promotion narrative.
- Confirm $|\text{corr vs momentum}| < 0.35$ for diversification claims.
- Drill report $\rightarrow$ `metrics_source_ids` $\rightarrow$ backtest payload before trusting a number.
- Read adversarial findings — HIGH/CRITICAL blocks promotion.
- Treat synthetic FX panels (`fx_synthetic_*`) as structural demos, not live edge.
- Run pytest invariants: `tests/test_quant_invariants.py`.

Extension rule: if a change lets natural language create a number that cannot be recomputed from stored returns + pinned formulas, reject the change.

检查 OOS 滚动验证夏普，而非仅样本内。

任何晋升叙述前在 PESSIMISTIC 成本下重跑。

对分散性声明确认 | 相对动量相关 | < 0.35 。

信任数字前先下钻 报告 $\rightarrow$ `metrics_source_ids` $\rightarrow$ 回测载荷。

阅读对抗性发现 — HIGH/CRITICAL 阻断晋升。

将合成 FX 面板 (`fx_synthetic_*`) 视为结构演示，非实盘优势。

运行 `pytest` 不变量: `tests/test_quant_invariants.py`。

扩展规则：若某改动使自然语言能产生无法由已存收益 + 固定公式重算的数字，则拒绝该改动。

27. Deep Dive: Walk-Forward Validation

27. 深入解析：滚动前瞻验证

File: engine/walk_forward.py. Walk-forward in QROS is **segmented out-of-sample evaluation of a frozen rule** — parameters in CrossSectionalConfig are treated as exogenous and are **not** re-fit inside each train window. This is intentionally stricter and more honest than nested parameter search that quietly overfits the train set.

文件: engine/walk_forward.py. QROS 中的 walk-forward 是对冻结规则的分段样本外评估——CrossSectionalConfig 中的参数视为外生，不会在每个训练窗内重新拟合。这比在训练集上静默过拟合的嵌套参数搜索更严格、更诚实。

27.1 Window modes and defaults

27.1 窗口模式与默认值

Field 字段	Default 默认	Meaning 含义
mode mode	expanding_window expanding_window	Train grows from t=0; alternate rolling_window 训练自 t=0 扩展; 可改 rolling
train_bars train_bars	126 126	~0.5y initial train (expanding) 扩展模式初始训练约半年
test_bars test_bars	63 63	~1 quarter OOS per window 每窗约一季样本外
step_bars step_bars	63 63	Advance train/test frontier 训练/测试前沿步进
rolling_train_bars rolling_train_bars	252 252	1y train when mode=rolling 滚动模式 1 年训练
cost_assumption cost_assumption	BASELINE BASELINE	Costs applied inside each slice backtest 每片回测内应用成本

27.2 Algorithm

27.2 算法

- Build window bounds (train_start, train_end, test_start, test_end) exclusive end indices.
- For each window, run full run_cross_sectional_backtest on the price slice (rolling mode prepends warm-up bars so lookback signals are defined).
- Extract only OOS returns in [test_start, test_end); store per-window metrics.
- Concatenate OOS segments with de-duplication so overlapping steps do not double-count.
- Aggregate OOS Sharpe / return / drawdown from the de-duplicated OOS path — this is the promotion-relevant number, not in-sample Sharpe.

```
# engine/walk_forward.py (conceptual)
for (ts, te, vs, ve) in _window_bounds(n, wf_cfg):
    slice_prices = prices.iloc[:ve] # expanding
    # rolling: include warm-up before ts
    bt = run_cross_sectional_backtest(slice_prices, cs_cfg, cost_model)
    oos = returns_in_original_index[vs:ve] # exclusive end
    windows.append(WalkForwardWindow(..., metrics=calculate_metrics(oos)))
aggregate = calculate_metrics(dedupe_concat(oos_parts))
```

构建窗口边界 (train_start...test_end)，终点索引为开区间。

每窗对价格切片跑完整

run_cross_sectional_backtest (滚动模式前置预热 K 线，使回看信号有定义)。

仅抽取 [test_start, test_end) 的样本外收益；保存每窗指标。

拼接样本外片段并去重，避免步进重叠导致双重计数。

由去重后的样本外路径汇总 OOS

夏普/收益/回撤——这才是与晋升相关的数字，而非样本内夏普。

Interpretation: a strategy with IS Sharpe 1.8 but aggregate OOS Sharpe ≤ 0 is a failed candidate — regardless of narrative quality.

解读：样本内夏普 1.8 但汇总 OOS 夏普 ≤ 0 的策略是失败候选——无论叙述多漂亮。




Figure 11. Expanding walk-forward train/test schedule.

图 11. 扩展式 walk-forward 训练/测试日程。

28. Deep Dive: Robustness & Fragility

28. 深入解析：稳健性与脆弱性

File: `engine/robustness.py` — `analyze_parameter_surface()`. The battery re-runs the cross-sectional backtest across a grid of one parameter (default: `lookback` $\in \{10,15,20,25,30,40,60\}$) and classifies the Sharpe surface.

文件: `engine/robustness.py` — `analyze_parameter_surface()`。套件在单一参数网格上重跑横截面回测（默认 `lookback` $\in \{10,15,20,25,30,40,60\}$ ），并对夏普曲面分类。

28.1 Fragility and smoothness rules

28.1 脆弱性与平滑规则

- **Fragile**: interior peak Sharpe exceeds neighbor mean by > 0.5 AND peak Sharpe $> 0.3 \rightarrow$ flagged as possible overfit knife-edge.
- **Boundary peak**: peak at first/last grid point $\rightarrow$ note to extend the grid before calling knife-edge.
- **Smooth**: mean $|\Delta \text{Sharpe}|$ across adjacent grid points < 0.35 .
- Preferred outcome: broad-ish stable region with positive Sharpe — not a single spike.

```
# Fragility test (interior peaks)
gap = peak_sharpe - mean(neighbor_sharpes)
fragile = (gap > 0.5) and (peak_sharpe > 0.3)
smooth = mean(abs(diff(sharpes))) < 0.35
```

脆弱：内部峰值夏普高于邻域均值 > 0.5 且峰值 $> 0.3 \rightarrow$ 标记为可能过拟合的刀锋峰。

边界峰：峰在网格首/末点 $\rightarrow$ 提示先扩展网格再下刀锋结论。

平滑：相邻网格点平均 $|\Delta \text{夏普}| < 0.35$ 。

更优结果：宽平台、正夏普区域——而非单点尖峰。




Figure 12. Illustrative smooth vs fragile lookback surfaces.

图 12. 平滑 vs 脆弱 lookback 曲面示意。

Adversarial review treats fragile=True / jagged surfaces as promotion blockers even when the chosen lookback prints a high in-sample Sharpe.

对抗性审查将 fragile=True / 锯齿曲面视为晋升阻断，即使所选 lookback 样本内夏普很高。

29. Deep Dive: Regime Analysis

29. 深入解析：市场状态分析

File: engine/regime.py. Labels are **heuristic** (rolling volatility tertiles × risk-on/off trend), not a claim of true latent states. Confidence is intentionally low (~0.4) so UI/report consumers do not over-trust the taxonomy.

- `vol = rolling_std(market_proxy, 20); trend = rolling_mean(market_proxy, 20).`
- `vol_bin ∈ {low_vol, mid_vol, high_vol}; trend_bin ∈ {risk_on, risk_off}.`
- Label string: "`{vol}|{trend}`"; shifted by 1 bar so day-t return is attributed to regimes known through t-1 (no same-day leakage).
- Per-regime metrics via pinned `calculate_metrics` (skip cells with < 5 observations).
- Concentration note when best regime Sharpe > 0.5 and worst < 0.

文件: engine/regime.py。标签是启发式的（滚动波动率三分位 × 风险偏好趋势），并非真实隐状态声明。置信度故意偏低（~0.4），避免 UI/报告读者过度信任该分类。

`vol` = 市场代理 20 日滚动标准差；`trend` = 20 日滚动均值。

`vol_bin ∈ {low_vol, mid_vol, high_vol}; trend_bin ∈ {risk_on, risk_off}.`

标签: "`{vol}|{trend}`"; 平移 1 根 K 线，使第 t 日收益归因于截至 t-1 已知状态（无同日泄漏）。

各状态指标经固定 `calculate_metrics` 计算（观测 < 5 的单元跳过）。

当最佳状态夏普 > 0.5 且最差 < 0 时附加集中度备注。

A strategy that only works in one historical regime slice is not book-ready — even if full-sample Sharpe looks fine.

仅在单一历史状态切片有效的策略不具备入册条件——即使全样本夏普看起来不错。

30. Deep Dive: Diversification Mathematics

30. 深入解析：分散化数学

File: alpha/registry.py — `analyze_diversification()`. Flagship workflow compares candidates against seeded ALP-existing-momentum.

文件: alpha/registry.py
`analyze_diversification()`。旗舰流程将候选与种子 ALP-existing-momentum 比较。

Output field 输出字段	Definition 定义	Promotion heuristic 晋升启发式
----------------------	------------------	------------------------------

return_correlations 收益相关	Pearson corr(candidate, existing_i) on aligned days 对齐日上 Pearson 相关	Track max corr vs book 跟踪相对组合的最大 相关
downside_correlations 下行相关	Corr on days where either series is negative 任一侧为负的日子上的相关	Crash co-movement matters more than calm corr 崩盘共动比平静相关更重要
genuine_diversification 真实分散	avg corr < 0.35 AND avg downside corr < 0.45 平均相关 < 0.35 且 平均下行相关 < 0.45	False → duplicate risk bet False → 重复风险押注
incremental_sharpe_hint 增量夏普提示	Equal-weight existing vs existing+candidate Sharpe (ddof=0) 既有等权 vs 既有+候选 夏普 (ddof=0)	delta > 0 preferred but not sufficient alone delta > 0 更优但单独不足

```
# Diversification core
corr = aligned.c.corr(aligned.e)
down_mask = (aligned.c < 0) | (aligned.e < 0)
down_corr = aligned.loc[down_mask].corr()
genuine = abs(avg_corr) < 0.35 and abs(avg_down) < 0.45
port_old = existing_df.mean(axis=1)
port_new = concat([port_old, candidate], axis=1).mean(axis=1)
delta = sharpe(port_new) - sharpe(port_old) # ddof=0 annualized
```

31. Deep Dive: ToolRouter & Allowlist

31. 深入解析：ToolRouter 与白名单

File: tools/router.py. Agents never call engine functions directly for research numbers. They invoke named tools through ToolRouter, which enforces an allowlist, records audit payloads, and returns structured ToolResult objects. This is the hard boundary that prevents LLM prose from inventing Sharpes.

文件: tools/router.py. 智能体从不直接调用引擎函数获取研究数字。它们经ToolRouter调用具名工具；路由强制白名单、记录审计载荷，并返回结构化ToolResult。这是阻止 LLM 叙述编造夏普的硬边界。

Tool 工具	Engine / store 引擎/存储	Produces 产出
create_experiment create_experiment	SQLite experiments SQLite experiments	experiment_id + configuration_hash 实验 ID + 配置哈希
run_backtest run_backtest	cross_sectional + metrics + costs 横截面+指标+成本	backtest_id, returns, PerformanceMetrics 回测 ID、收益、指标
walk_forward walk_forward	walk_forward.py walk_forward.py	OOS windows + aggregate 样本外窗 + 汇总
robustness_battery robustness_battery	robustness.py robustness.py	surface, fragile, smooth 曲面、脆弱、平滑
analyze_diversification analyze_diversification	alpha/registry.py alpha/registry.py	corr, genuine_diversification 相关、真实分散
run_stress_test run_stress_test	portfolio.py portfolio.py	stress scenarios 压力情景
run_bootstrap run_bootstrap	statistics.py statistics.py	uncertainty / multiple-testing aware stats 不确定性/多重检验感知统计

ToolRouter: agents call tools; engines compute numbers




Figure 13. Numbers flow Agent → ToolRouter → Engine → Storage.

图 13. 数字流：智能体 → ToolRouter → 引擎 → 存储。

Extension rule: new research capabilities must land as allowlisted tools with deterministic engines behind them — never as free-form model arithmetic.

扩展规则：新研究能力必须以白名单工具落地、背后是确定性引擎——绝非自由形式的模型算术。

32. Worked Example: From Returns to Decision

32. 演算示例：从收益到决策

The following toy path illustrates how a researcher should mentally audit a candidate. Numbers are pedagogical (not a live run dump).

以下玩具路径说明研究员应如何心智审计一个候选。数字用于教学（非实况运行转储）。

32.1 Net daily returns → pinned metrics

32.1 净日收益 → 固定指标

Suppose a 5-day net return series after costs: $r = [+0.4\%, -0.2\%, +0.3\%, -0.1\%, +0.2\%]$. Equity path: $1.004 \rightarrow 1.0020 \rightarrow 1.0050 \rightarrow 1.0040 \rightarrow 1.0060$. Cumulative $\approx +0.60\%$. With $N=5$ the annualization factor is noisy — in production N is hundreds of bars — but the formulas stay identical: $\sigma = \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$, Sharpe = $\text{mean}(r) / \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$ ($r_f=0$).

假设成本后 5 日净收益序列: $r = [+0.4\%, -0.2\%, +0.3\%, -0.1\%, +0.2\%]$ 。权益路径: $1.004 \rightarrow 1.0020 \rightarrow 1.0050 \rightarrow 1.0040 \rightarrow 1.0060$ 。累计 $\approx +0.60\%$ 。N=5 时年化噪声大——生产中 N 为数百根 K 线——但公式不变: $\sigma = \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$, Sharpe = $\text{mean}(r) / \text{std}(r, \text{ddof}=0) \cdot \sqrt{252}$ ($r_f=0$)。

32.2 Cost stress

32.2 成本压力

If average one-way turnover per rebalance is 0.40 and `rebalance_every=5`, daily average turnover $\approx 0.40/5 = 0.08$. Under BASELINE (2.75 bps variable) daily cost $\approx 0.08 \times 2.75e-4 = 2.2e-5$ (~0.22 bps/day). Under PESSIMISTIC (10 bps) ≈ 0.8 bps/day. A thin edge that disappears under this haircut is not promotable.

若每次再平衡单边换手 0.40 且 `rebalance_every=5`, 日均换手 $\approx 0.40/5 = 0.08$ 。BASELINE (2.75 bps 可变) 下日成本 $\approx 0.08 \times 2.75e-4 = 2.2e-5$ (约 0.22 bps/日)。PESSIMISTIC (10 bps) 下 ≈ 0.8 bps/日。被这层摩擦吃掉的薄优势不可晋升。

32.3 Decision table

32.3 决策表

Check 检查	Pass criterion 通过标准	Example outcome 示例结果
IS Sharpe (audit only) 样本内夏普 (仅审计)	Recomputable from returns 可由收益重算	1.1 — not decisive 1.1 — 非决定性
OOS Sharpe OOS 夏普	> 0.3 after BASELINE costs BASELINE 成本后 > 0.3	0.18 — fail 0.18 — 失败
Fragility 脆弱性	fragile=False, smooth=True fragile=False, smooth≈True	fragile=True — fail fragile=True — 失败
corr(momentum) corr(动量)	< 0.35 < 0.35	0.62 — duplicate risk 0.62 — 重复风险
Adversarial 对抗性	severity < CRITICAL 严重度 < CRITICAL	HIGH cost sensitivity HIGH 成本敏感
Decision 决策	PROMOTABLE / MORE_RESEARCH / REJECT 可晋升 / 需更多研究 / 拒绝	REQUIRES_MORE_RESEARCH REQUIRES_MORE_RESEARCH
Honest outcome for the flagship synthetic FX demo is often REQUIRES_MORE_RESEARCH — that is a feature of scientific hygiene, not a product defect.		旗舰合成外汇演示的诚实结果常常是 REQUIRES_MORE_RESEARCH——这是科研卫生的特性，不是产品缺陷。

33. Workstation Page-by-Page Guide

33. 工作站逐页指南

The Next.js app under `web/` is the daily operator surface. Each page has one primary job. Use this map when onboarding or auditing a run.

`web/` 下的 Next.js 应用是日常操作面。每页只有一个主职责。入门或审计运行时用此地图。

Route 路由	Primary job 主职责	Evidence to open next 下一步打开的证据
/	Fleet pulse: latest research, counts, health 总览: 最新研究、计数、健康	Click research → /research/[id] 点研究 → /research/[id]

/research /research	Queue of research requests + statuses 研究请求队列与状态	Open workspace for a run 打开某次运行的工作区
/research/[id] /research/[id]	Workflow graph, plan, agents, report 工作流图、计划、智能体、报告	Experiments sorted by Sharpe 按夏普排序的实验
/experiments /experiments	All experiments with enriched metrics 全部实验（富化指标）	Row → config hash + backtest 行 → 配置哈希 + 回测
/alphas /alphas	Library status + metrics_source_ids 库状态 + metrics_source_ids	Drill to source backtest 下钻到来源回测
/portfolio /portfolio	What-if book + Alpha X 情景：组合 + Alpha X	Risk page for concentration 风险页看集中度
/risk /risk	Correlation / concentration alerts 相关/集中度告警	Reject high-corr adds 拒绝高相关新增
/data /data	Dataset catalog / synthetic panels 数据集目录/合成面板	Confirm fx_synthetic_* labels 确认 fx_synthetic_* 标签
/regimes /regimes	Regime performance breakdown 状态绩效分解	Check concentration notes 检查集中度备注
/paper /paper	Paper mode (IID honesty banners) 模拟盘（IID 诚实横幅）	Never treat as LIVE 绝不当作实盘
/reports /reports	Decision + survivors + narrative 决策 + 幸存者 + 叙述	Traceability chain 可追溯链
/agents /agents	Allowlist / agent roles 白名单/角色	Inspector payloads 检查器载荷
/memory /memory	Prior findings / memory store 既有发现/记忆库	Avoid repeating failed ideas 避免重复失败想法
/system /system	API health, versions, ops API 健康、版本、运维	Restart / key config 重启/密钥配置

33.1 Research workspace mental model

33.1 研究工作区心智模型

- Graph nodes = orchestrator stages; green/completed means tools finished, not that alpha is good.
 - Agent Inspector shows tool args/results — this is the audit log for numbers.
 - Report decision is downstream of validation; always verify survivors' metrics_source_ids.
 - Bottom command bar starts new research with budget (max_experiments) — budgets are first-class.
- 图节点 = 编排阶段；绿色/完成表示工具跑完，不表示 Alpha 优良。

智能体检查器显示工具参数/结果——这是数字的审计日志。

报告决策在验证下游；务必核幸存者 metrics_source_ids。

底部命令栏以预算（max_experiments）启动新研究——预算是一等公民。

34. Translation & Review Notes

34. 翻译与审阅说明

This section records terminology choices and a reviewer checklist for the bilingual edition.

本节记录双语版术语选择与审阅清单。

34.1 Terminology glossary

34.1 术语表

English 英文	Chinese (chosen) 中文（选用）	Notes 说明
Alpha Alpha	Alpha■■■■■■■■■ Alpha	Keep “Alpha” as loanword 保留英文借词
Sharpe ratio Sharpe	■■■■■ 夏普比率	Standard CN finance term 标准金融译名
Drawdown Drawdown	■■■ / ■■■■■ 回撤	Max DD → ■■■■■ 最大回撤
Walk-forward Walk-forward	■■■■■■■ 滚动前验验证	Also ■■■■■ 亦可称走走验证
Look-ahead bias Look-ahead	■■■■■ 前视偏差	Decision-time info only 仅决策时可得信息

Cross-sectional CS	████ 横截面	FX basket ranking 外汇篮子排序
Adversarial review Adversarial	██████ 对抗性审查	Critical review, not GAN 批评式审查，非生成模型
Orchestrator Orchestrator	████ 编排器	Research state machine 研究状态机
Deterministic Deterministic	████ 确定性	Opposite of stochastic LLM 相对随机 LLM
Paper trading Paper	██████ 模拟盘	Labeled simulation mode 已标注的仿真模式
Backtest Backtest	███ 回测	Standard 标准译名
Regime Regime	█████ 市场状态	Also ███ 亦可称区制
Turnover Turnover	████ 换手率	Weight-space turnover 权重空间换手
Provenance Provenance	███ / █████ 溯源	metrics_source_ids 指标来源 ID
Allowlist Allowlist	████ 白名单	Tool permit-list 工具许可名单

34.2 Reviewer checklist (this edition)

34.2 本版审阅清单

- No promotional exaggeration: paper mode and deterministic agents are disclosed in both languages.
- No inappropriate content; tone is institutional and technical.
- Quant terms use standard Mainland Chinese finance translations; English identifiers preserved.
- “Adversarial” means critical review of strategies (██████), not generative-model jargon.
- Figures, captions, and tables are bilingual; code remains English with optional ZH comments.
- Side-by-side layout: English left column, Chinese right column with light background.

无夸大宣传：模拟盘模式与确定性智能体在两种语言中均如实披露。

无不当内容；语气保持机构化与技术性。

量化术语采用大陆常用金融译名；英文标识符保留。

“Adversarial”
指对策略的批评式审查（对抗性审查），而非生成模型术语。

图、图注与表格双语；代码保持英文并可选中文注释。

对照版式：左英右中，中文列浅底以利扫读。

End of bilingual handbook. Optimize for researchers using the system for hours every day: traceability over spectacle, engines over narratives, budgets over unbounded loops.

双语手册完。为每日长时间使用的研究员优化：可追溯优于炫技，引擎优于叙述，预算优于无界循环。